


Concurrency in Python

An overview of multithreading and multiprocessing

Why?



But remember...

Premature optimization is the
root of all evil.

-Donald Knuth

Problems Concurrency Solves

I/O bound (multithreading & asynchronous)

- Your program is spending more time than you want sending or receiving data.
- Examples: web requests, harddrive i/o, slow API calls, etc.

CPU bound (multiprocessing)


- Your program is spending more time than you want pegging the processor.
- Examples: calculations, image processing, data analysis, etc.

Or both...

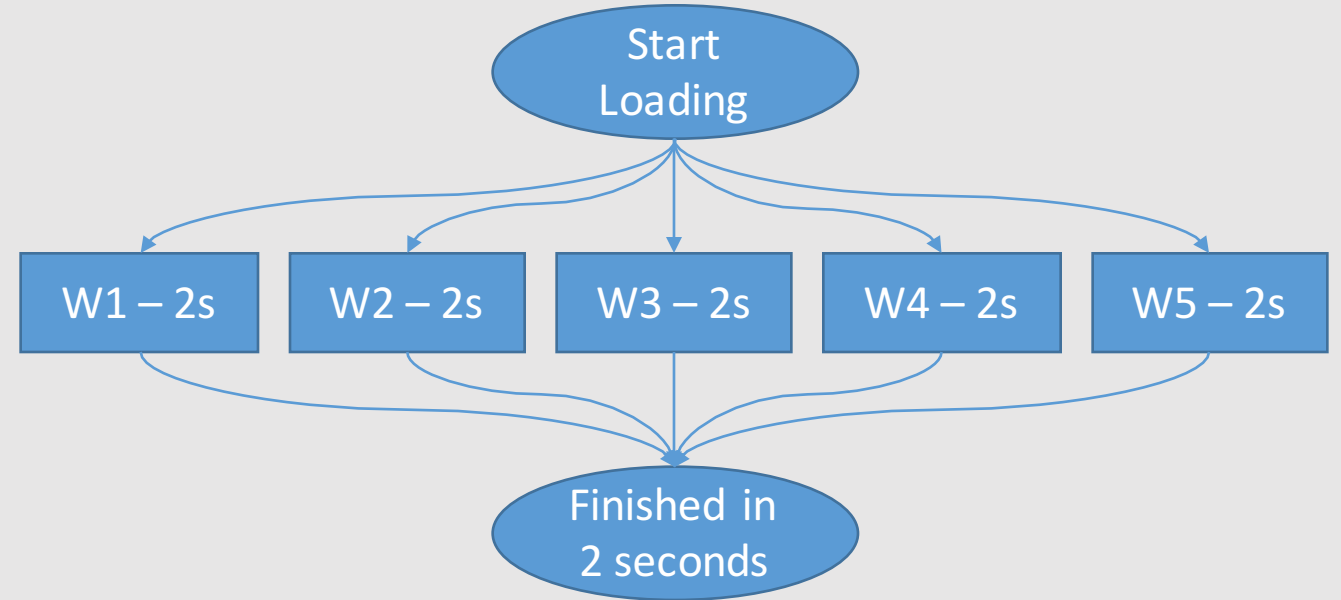
- E.g. loading and analyzing web pages or other data.

Using Multithreading

Simple procedural way:




Multithreaded with 5 threads:

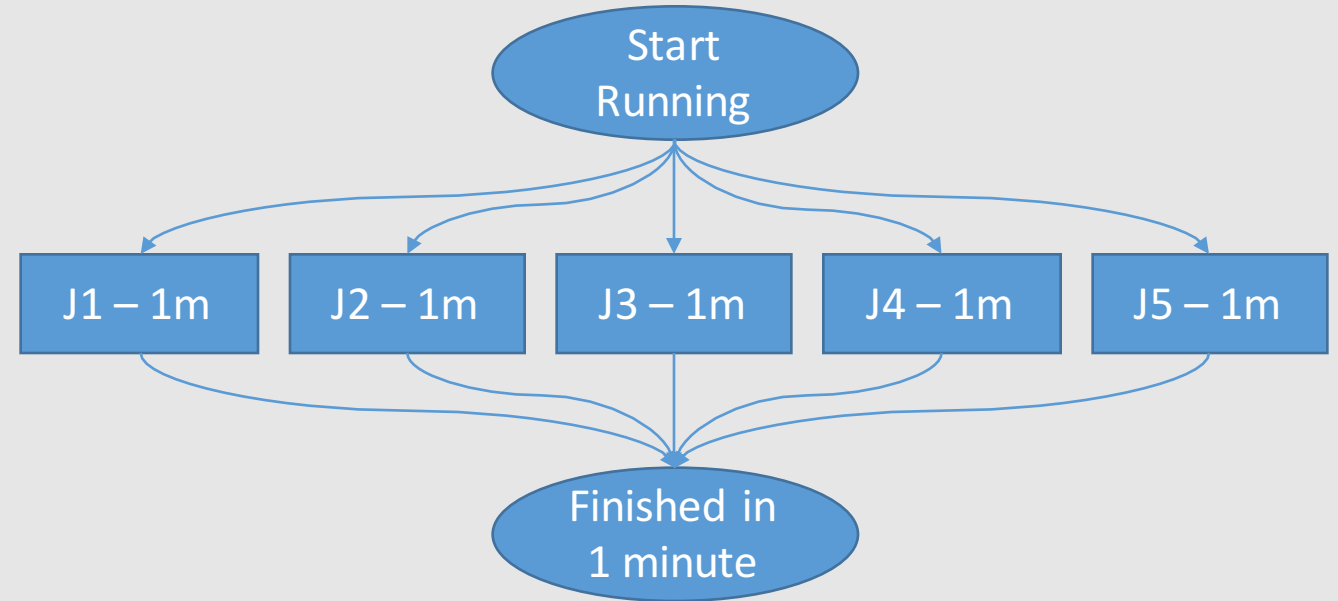


Using Multiprocessing

Simple procedural way:



Multiprocessed with 5 processes and 5 available cores:



Details...

Multithreading

- Threads share global memory.
- Threads have independent stacks.
- Because of the GIL, execution is only logically parallel, not actually.
- Threads can communicate through a variety of methods since they share global memory.

Multiprocessing

- Execution paths are completely independent (they do not share global memory)
- Allows you to use multiple cores and actually run in parallel.
- Communication between processes is more restricted since memory is not shared (pipes, sockets, message queues, explicitly shared memory).

A Worker Queue Example

- In Python multithreading and multiprocessing programs can be written very similarly.
- But there are some important differences...

Multithreading

```
1 import threading as tr
2 import Queue
3 import time
4
5 globalID = None
6 start = time.time()
7
8 # Function to call on every job
9 def jobThread(jobQueue, threadID):
10     global globalID
11     while True:
12         try:
13             job = jobQueue.get_nowait()
14         except Queue.Empty:
15             break
16         else:
17             sleeptime = job
18             now = time.time() - start
19             print("{:.2f}: [{}] read globalID {}, setting to {}".format(
20                 now, threadID, globalID, threadID))
21             globalID = threadID
22             time.sleep(sleeptime)
23             jobQueue.task_done()
24
25 # Function to start the threads
26 def doWork(jobs, maxworkers=10):
27     jobCount = len(jobs)
28     jobQueue = Queue.Queue()
29     for job in jobs:
30         jobQueue.put(job)
31
32     maxworkers = min(maxworkers, jobCount)
33
34     threads = [tr.Thread(target=jobThread, args=(jobQueue,n)) for n in range(maxworkers)]
35     for t in threads:
36         t.start()
37
38     jobQueue.join()
39
40 # create job list
41 jobList = []
42 for i in xrange(20):
43     jobList.append(1)
44
45 doWork(jobList)
46 print("{:.2f}: All done!".format(time.time() - start))
```

SIMPLE EXAMPLES

Multiprocessing

```
1 import multiprocessing as mp
2 import Queue
3 import time
4
5 globalID = None
6 start = time.time()
7
8 # Function to call on every job
9 def jobThread(jobQueue, threadID):
10     global globalID
11     while True:
12         try:
13             job = jobQueue.get_nowait()
14         except Queue.Empty:
15             if jobQueue.empty(): break # get_nowait() can timeout on a non-empty queue
16         else:
17             sleeptime = job
18             now = time.time() - start
19             print("{:.2f}: [{}] read globalID={}, setting to {}".format(
20                 now, threadID, globalID, threadID))
21             globalID = threadID
22             time.sleep(sleeptime)
23             jobQueue.task_done()
24
25 # Function to start the processes
26 def doWork(jobs, maxworkers=10):
27     jobCount = len(jobs)
28     jobQueue = mp.JoinableQueue()
29     for job in jobs:
30         jobQueue.put(job)
31
32     maxworkers = min(maxworkers, jobCount)
33
34     processes = [mp.Process(target=jobThread, args=(jobQueue,n)) for n in range(maxworkers)]
35     for p in processes:
36         p.start()
37
38     jobQueue.join()
39
40 # create job list
41 jobList = []
42 for i in xrange(20):
43     jobList.append(1)
44
45 doWork(jobList)
46 print("{:.2f}: All done!".format(time.time() - start))
```

```

thread_example$ python difference_thread.py
0.00: [0] read globalID None, setting to 0
0.00: [1] read globalID 0, setting to 1
0.00: [2] read globalID 1, setting to 2
0.00: [3] read globalID 2, setting to 3
0.00: [4] read globalID 3, setting to 4
0.00: [5] read globalID 3, setting to 5
0.00: [6] read globalID 3, setting to 6
0.00: [7] read globalID 5, setting to 7
0.00: [8] read globalID 7, setting to 8
0.00: [9] read globalID 8, setting to 9
1.00: [0] read globalID 9, setting to 0
1.00: [8] read globalID 0, setting to 8
1.00: [9] read globalID 0, setting to 9
1.00: [2] read globalID 9, setting to 2
1.00: [1] read globalID 2, setting to 1
1.00: [4] read globalID 1, setting to 4
1.00: [6] read globalID 1, setting to 6
1.00: [5] read globalID 6, setting to 5
1.00: [7] read globalID 5, setting to 7
1.00: [3] read globalID 5, setting to 3
2.00: All done!
thread_example$

```

Contention / Races

Multithreading

SIMPLE EXAMPLES OUTPUT

```

thread_example$ python difference_process.py
0.05: [0] read globalID=None, setting to 0
0.05: [1] read globalID=None, setting to 1
0.05: [2] read globalID=None, setting to 2
0.05: [3] read globalID=None, setting to 3
0.05: [4] read globalID=None, setting to 4
0.05: [5] read globalID=None, setting to 5
0.05: [6] read globalID=None, setting to 6
0.05: [7] read globalID=None, setting to 7
0.05: [8] read globalID=None, setting to 8
0.05: [9] read globalID=None, setting to 9
1.05: [0] read globalID=0, setting to 0
1.05: [1] read globalID=1, setting to 1
1.05: [2] read globalID=2, setting to 2
1.05: [3] read globalID=3, setting to 3
1.05: [4] read globalID=4, setting to 4
1.05: [5] read globalID=5, setting to 5
1.05: [6] read globalID=6, setting to 6
1.05: [7] read globalID=7, setting to 7
1.05: [8] read globalID=8, setting to 8
1.05: [9] read globalID=9, setting to 9
2.06: All done!
thread_example$

```

Multiprocessing

Multithreading

```
1 import threading as tr
2 import Queue
3 import time
4
5 globalID = None
6 start = time.time()
7
8 # Function to call on every job
9 def jobThread(jobQueue, threadID):
10     global globalID
11     while True:
12         try:
13             job = jobQueue.get_nowait()
14         except Queue.Empty:
15             break
16         else:
17             sleep(1)
18             now = time.time()
19             print(f"Thread {threadID} processing job {job} at {now}")
20             globalID = now
21             jobQueue.put(job)
22             print(f"Thread {threadID} finished job {job} at {now}")
23             jobQueue.get_nowait()
24             print(f"Thread {threadID} finished job {job} at {now}")
25
26 # Function to call on every job
27 def doWork(jobList):
28     jobCount = len(jobList)
29     jobQueue = Queue.Queue()
30     for job in jobList:
31         jobQueue.put(job)
32
33     maxworkers = min(maxworkers, jobCount)
34
35     threads = [tr.Thread(target=jobThread, args=(jobQueue,n)) for n in range(1, maxworkers+1)]
36     for t in threads:
37         t.start()
38
39     jobQueue.join()
40
41 # create job list
42 jobList = []
43 for i in xrange(20):
44     jobList.append(1)
45
46 doWork(jobList)
47 print("{:.2f}: All done!".format(time.time() - start))
```

Thread #0

Thread #1

Thread #2

Thread #3

Thread #4

Thread #5

Thread #6

Thread #7

Thread #8

Thread #9



Problems Concurrency Creates

- Adds complexity:
 - Can be much more difficult to debug.
 - Can have unknown bugs and race conditions.
- Switching between multithreading and multiprocessing can be confusing b/c of the different memory scopes.
- Proper testing is more difficult (edge cases, etc) and even impractical.
- Can hurt performance (e.g. multithreading CPU bound code).
- Exceptions can disappear if you don't catch them.
- Program may not die normally, may need a SIGTERM or SIGKILL.

How To...

1. Write a mostly usable modular program without concurrency.
 2. Find the bottlenecks in performance (I/O or CPU bound).
 3. Fix the worst bottlenecks by adding concurrency where appropriate.
 4. Repeat from step 2 until you are satisfied with the performance or there is nothing left to optimize (within reason).
- Avoid optimizing things that don't need to be optimized (see Knuth).

MOAR

Details...

Communication

- Multithreading:
 - Global variables (not ideal)
 - Shared variables (i.e. thread function parameters)
 - Queues
- Multiprocessing:
 - Multiprocessing.Queue
 - Multiprocessing.Pipe
 - UNIX sockets or TCP sockets
 - Shared memory (memory mapped files)
 - Windows and Linux message queues (OS persistence)

Locks

- Safely control access to global or shared variables.
- Use 'with'!
- If a module or class is 'threadsafe', it does that for you (like Queue).

```
1 import threading as tr
2 import time
3
4 globalDict = {} # Be very careful with global mutable objects
5 globalDictLock = tr.RLock() # Just use RLock, don't deadlock yourself
6 shutdown = False
7
8 def taskThread1(taskparams):
9     while not shutdown:
10         # ... do task ...
11         state = 'state1'
12         with globalDictLock:
13             globalDict['task1'] = state
14             time.sleep(0.1)
15
16
17 def taskThread2(taskparams):
18     while not shutdown:
19         # ... do task ...
20         state = 'state2'
21         with globalDictLock:
22             globalDict['task2'] = state
23             time.sleep(0.1)
```

Termination Requests

- Python doesn't kill children
- Use signal for clean exits

```
1 import signal
2 import time
3
4 shutdown = False
5
6 def sighandler(signum, frame):
7     global shutdown
8     print("Detected a shutdown request with signal {}".format(signum))
9     shutdown = True
10
11 signal.signal(signal.SIGTERM, sighandler)
12 signal.signal(signal.SIGINT, sighandler)
13 signal.signal(signal.SIGHUP, sighandler)
14
15 while not shutdown:
16     time.sleep(0.1)
```

External Modules

- Can be helpful but can also add more unnecessary complexity.
- E.g. I tried using grequests, but for my use case it was much more effort, adding numerous callbacks and conditions... A simple threaded worker function gave me total control of the threads.

Resources

- The Little Book of Semaphores (free) - <http://www.greenteapress.com/semaphores/>
- <https://docs.python.org/2/library/threading.html> (or 3.5)
- <https://docs.python.org/2/library/multiprocessing.html> (or 3.5)